


Assignment - 3

Team 3

April 13, 2025



CS 331

Computer Networks

Prof. Sameer G Kulkarni

INDIAN INSTITUTE OF TECHNOLOGY GANDHINAGAR
Palaj, Gandhinagar - 382355

Network Loops, NAT, and Routing Algorithms

[GitHub Repo](#)

Submitted by

Guntas Singh Saran (22110089)

Computer Science and Engineering

Hitesh Kumar (22110098)

Computer Science and Engineering

Contents

1	Network Loop	1
1.1	Creating the Topology	1
1.2	Network Behavior on Ping Tests	2
1.3	Solution: Spanning Tree Protocol (STP)	4
1.4	Ping Test Results	5
1.5	Conclusion	5
2	Configure Host-based NAT	7
2.1	Topology and Changes	8
2.2	Observations	8
2.2.1	Parts A and B Communication	8
2.2.2	H1 and H2 Communication	9
2.3	Issues and Rectifications	9
2.3.1	Initial Issues	9
2.3.2	Rectifications with Code	10
2.4	Finally all the Communication Tests	15
2.4.1	Communication to an external host from an internal host	15
2.4.2	Test communication to an internal host from an external host	15
2.4.3	iperf tests: 3 tests of 120s each.	15
3	Distributed Asynchronous Distance Vector Routing	18
3.1	Distance Vector Routing Algorithm	18
3.2	Implementation Approach	18
3.2.1	Data Structures	18
3.2.2	Code Organization	19
3.2.3	Key Logic	19
3.2.4	Refinements	19
3.3	Role of <code>distance_vector.c</code>	20

Chapter 1

Network Loop

1.1 Creating the Topology

The network topology comprises four switches (**s1** to **s4**) and eight hosts (**h1** to **h8**). Each host is connected to a switch and configured with a unique IP address within the same subnet. The IP configuration for each host can be easily understood from the given python snippet for the custom Topology Class.

Listing 1.1: Sample Python Listing `main.py`

```
1 class CustomTopo(Topo):
2     def build(self):
3         # Switches
4         s1 = self.addSwitch('s1')
5         s2 = self.addSwitch('s2')
6         s3 = self.addSwitch('s3')
7         s4 = self.addSwitch('s4')
8
9         h1 = self.addHost('h1', ip='10.0.0.2/24')
10        h2 = self.addHost('h2', ip='10.0.0.3/24')
11        h3 = self.addHost('h3', ip='10.0.0.4/24')
12        h4 = self.addHost('h4', ip='10.0.0.5/24')
13        h5 = self.addHost('h5', ip='10.0.0.6/24')
14        h6 = self.addHost('h6', ip='10.0.0.7/24')
15        h7 = self.addHost('h7', ip='10.0.0.8/24')
16        h8 = self.addHost('h8', ip='10.0.0.9/24')
17
18        # Host - Switch with 5ms delay
19        self.addLink(h1, s1, delay='5ms')
20        self.addLink(h2, s1, delay='5ms')
21        self.addLink(h3, s2, delay='5ms')
22        self.addLink(h4, s2, delay='5ms')
23        self.addLink(h5, s3, delay='5ms')
24        self.addLink(h6, s3, delay='5ms')
25        self.addLink(h7, s4, delay='5ms')
26        self.addLink(h8, s4, delay='5ms')
27
28        # Switch - Switch with 7ms delay
29        self.addLink(s1, s2, delay='7ms')
30        self.addLink(s2, s3, delay='7ms')
31        self.addLink(s3, s4, delay='7ms')
32        self.addLink(s4, s1, delay='7ms')
33        self.addLink(s1, s3, delay='7ms')
```

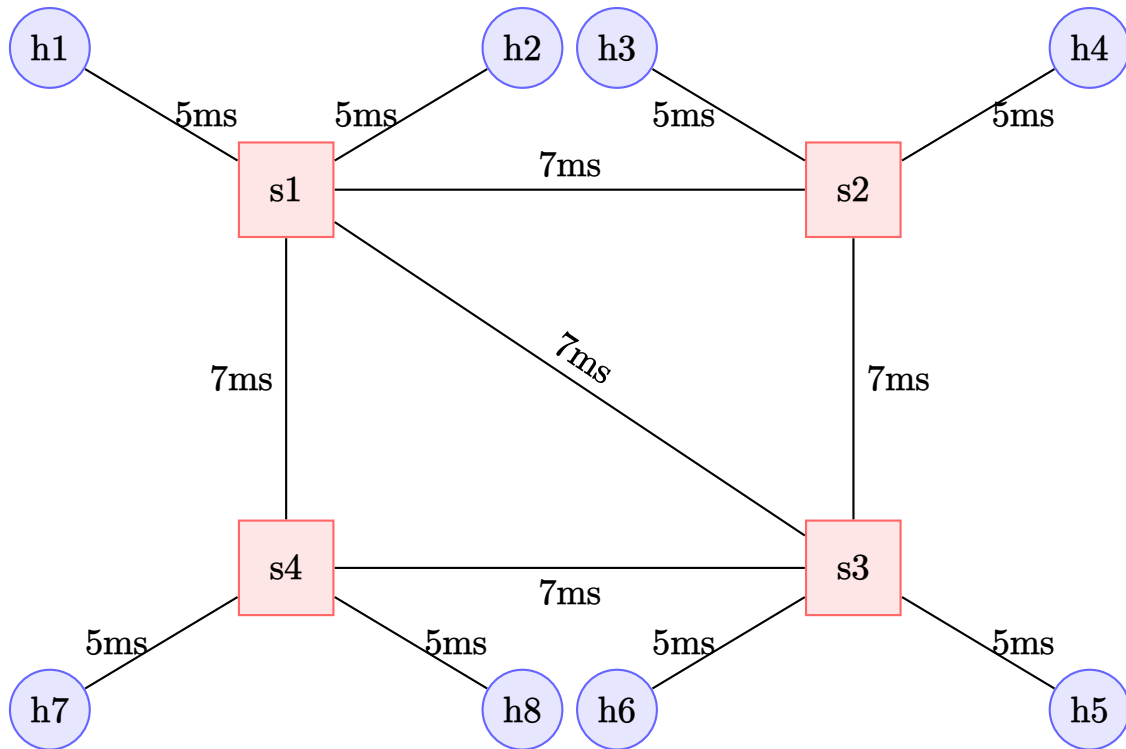


Figure 1.1: Network Topology with Four Switches and Eight Hosts

1.2 Network Behavior on Ping Tests

To evaluate the connectivity and performance of the network, a set of ICMP ping tests were conducted between selected host pairs. Each test was executed three times with a 30-second interval to allow for accurate measurement of latency and to observe consistency.

Ping Tests Performed

- h3 → h1
- h5 → h7
- h8 → h2

Observations

During the initial ping tests, all packets failed to reach their destinations. Each ping resulted in 100% packet loss. This indicated that there was a problem in the network's ability to forward packets correctly between switches.

Analysis

Upon investigation, it was found that the failure was due to the presence of loops in the switch topology. While having redundant paths between switches is generally beneficial for fault tolerance, it can lead to serious network issues if not managed properly.

Switches forward broadcast frames (such as ARP requests) out of all ports except the one they were received on. When there are loops, these broadcast frames keep circulating endlessly through the network. This phenomenon is known as a **broadcast storm**, and it consumes significant bandwidth, leading to degraded performance or complete network failure.



Additionally, Ethernet switches maintain a MAC address table to associate devices with specific ports. In a looped topology, the same frame may arrive at a switch from multiple directions, causing it to constantly update its MAC table with conflicting port information. This leads to **MAC table instability**, where the switch can no longer correctly determine which port to use for a particular destination. As a result, unicast traffic is flooded across the network just like broadcasts, worsening congestion.

Packet Capture Evidence

Packet captures during the ping tests confirmed that ARP requests were being broadcast repeatedly, and no ARP replies were received initially. This behavior is consistent with a switching loop, where broadcast frames circulate endlessly, preventing any meaningful traffic from being delivered.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
2	0.044332369	2e:eb:67:61:27:41	92:57:d9:f4:05:69	ARP		42 10.0.0.2 is at 2e:eb:67:61:27:41
4	0.046193520	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
5	0.046350525	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
6	0.046474564	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
7	0.046599442	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
8	0.047240592	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
9	0.047845424	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
10	0.081102802	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
11	0.081322944	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
12	0.081449777	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
13	0.081694783	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
14	0.081820778	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
15	0.081938951	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
16	0.086589871	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
17	0.086760007	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4

Figure 1.2: First reply to the ARP request from h1 to h3 with h1's MAC Address

No.	Time	Source	Destination	Protocol	Length	Info
43078	80.502314710	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43077	80.502198773	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43076	80.502075851	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43075	80.501954884	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43074	80.501834756	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43073	80.501701777	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43072	80.499761845	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43071	80.499646745	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43070	80.499529690	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43069	80.499412634	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43068	80.499294741	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43067	80.499176848	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43066	80.499059793	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43065	80.498942737	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43064	80.498825682	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
43063	80.498707789	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4

Figure 1.3: Flooding behavior observed as the same ARP packet is repeatedly rebroadcast

No.	Time	Source	Destination	Protocol	Length	Info
18281	19.094436231	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
18282	19.094464448	2e:eb:67:61:27:41	92:57:d9:f4:05:69	ARP		42 10.0.0.2 is at 2e:eb:67:61:27:41
18283	19.094576479	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
18284	19.094604417	2e:eb:67:61:27:41	92:57:d9:f4:05:69	ARP		42 10.0.0.2 is at 2e:eb:67:61:27:41
18285	19.094716447	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
18286	19.094745223	2e:eb:67:61:27:41	92:57:d9:f4:05:69	ARP		42 10.0.0.2 is at 2e:eb:67:61:27:41
18287	19.094857254	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
18288	19.094885471	2e:eb:67:61:27:41	92:57:d9:f4:05:69	ARP		42 10.0.0.2 is at 2e:eb:67:61:27:41
18289	19.094997502	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
18290	19.095026278	2e:eb:67:61:27:41	92:57:d9:f4:05:69	ARP		42 10.0.0.2 is at 2e:eb:67:61:27:41
18291	19.095138308	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
18292	19.095167364	2e:eb:67:61:27:41	92:57:d9:f4:05:69	ARP		42 10.0.0.2 is at 2e:eb:67:61:27:41
18293	19.099686584	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
18294	19.099767604	2e:eb:67:61:27:41	92:57:d9:f4:05:69	ARP		42 10.0.0.2 is at 2e:eb:67:61:27:41
18295	19.099967638	92:57:d9:f4:05:69	Broadcast	ARP		42 Who has 10.0.0.2? Tell 10.0.0.4
18296	19.100006605	2e:eb:67:61:27:41	92:57:d9:f4:05:69	ARP		42 10.0.0.2 is at 2e:eb:67:61:27:41

Figure 1.4: Response from h1 after STP was enabled and loops were eliminated



1.3 Solution: Spanning Tree Protocol (STP)

Spanning Tree Protocol (STP) is a Layer 2 protocol that prevents loops in Ethernet networks by selectively blocking redundant paths while keeping one active path between any two nodes. In our initial configuration, STP was not enabled, which led to broadcast storms and MAC table instability due to the presence of loops in the network. To resolve this, STP was enabled on all switches to ensure a loop-free topology.

Once STP was activated, the resulting topology and port states were examined using Mininet commands. The first observation was the full connectivity of hosts and switches as shown in the output of the `net` command below:

```
1 mininet> net
2 h1 h1-eth0:s1-eth1
3 h2 h2-eth0:s1-eth2
4 h3 h3-eth0:s2-eth1
5 h4 h4-eth0:s2-eth2
6 h5 h5-eth0:s3-eth1
7 h6 h6-eth0:s3-eth2
8 h7 h7-eth0:s4-eth1
9 h8 h8-eth0:s4-eth2
10 s1 lo: s1-eth1:h1-eth0 s1-eth2:h2-eth0 s1-eth3:s2-eth3 s1-eth4:s4-eth4
    s1-eth5:s3-eth5
11 s2 lo: s2-eth1:h3-eth0 s2-eth2:h4-eth0 s2-eth3:s1-eth3 s2-eth4:s3-eth3
12 s3 lo: s3-eth1:h5-eth0 s3-eth2:h6-eth0 s3-eth3:s2-eth4 s3-eth4:s4-eth3
    s3-eth5:s1-eth5
13 s4 lo: s4-eth1:h7-eth0 s4-eth2:h8-eth0 s4-eth3:s3-eth4 s4-eth4:s1-eth4
14 c0
```

Next, the status of each switch interface was inspected using the `dpctl show` command. From the output, it was evident that STP had successfully blocked the redundant paths by placing the following ports into `STP_BLOCK` state:

- `s2-eth4`: s2's interface connected with s3 switch's `s3-eth3` interface.
- `s3-eth5`: s3's interface connected with s1 switch's `s1-eth5` interface.

NOTE: These blocked interfaces might differ each time we create the topology in mininet and start STP. So, the results for another run may differ.

All remaining interfaces were in the `STP_FORWARD` state, actively handling traffic while preventing any loops in the network. But, STP takes some time to converge and during that time the state of all those interfaces remains `STP_LEARN`.

```
3(s4-eth3): addr:ca:fc:13:15:d0:a6
config:      0
state:       STP_LEARN
current:     10GB-FD COPPER
speed: 10000 Mbps now, 0 Mbps max
```

Figure 1.5: STP Learn state while the Redundant links are to be found

To visually represent the updated, loop-free topology after STP convergence, the TikZ diagram shown below highlights only the active links. Blocked interfaces (s2-eth4 and s3-eth5) have been excluded to reflect the logical topology used for data forwarding:




Figure 1.6: Effective Network Topology after STP has Blocked Redundant Interfaces

1.4 Ping Test Results

Table 1.1: Average RTTs for 30s Ping Tests After STP Convergence

Ping Test	Attempt 1 (ms)	Attempt 2 (ms)	Attempt 3 (ms)	Avg. Time (ms)
h3 → h1	34.469	34.327	34.763	34.520
h5 → h7	35.326	34.440	34.248	34.671
h2 → h8	35.934	34.731	34.818	35.161

The ping results observed for all the three ping configurations can also be inferred from the fig. 1.6. To show the correctness of the Network Topology let's try to ping h5 from h4 and for that the expected RTT is about 62 ms through $s2 \rightarrow s1 \rightarrow s4 \rightarrow s3$. The figure below shows the similar results with RTT having 65 ms value nearly equal to 62 ms.

```
--- 10.0.0.6 ping statistics ---
30 packets transmitted, 30 received, 0% packet loss, time 29037ms
rtt min/avg/max/mdev = 62.125/65.054/130.920/12.235 ms
```

Figure 1.7: Ping Result for $h4 \rightarrow h5$

1.5 Conclusion

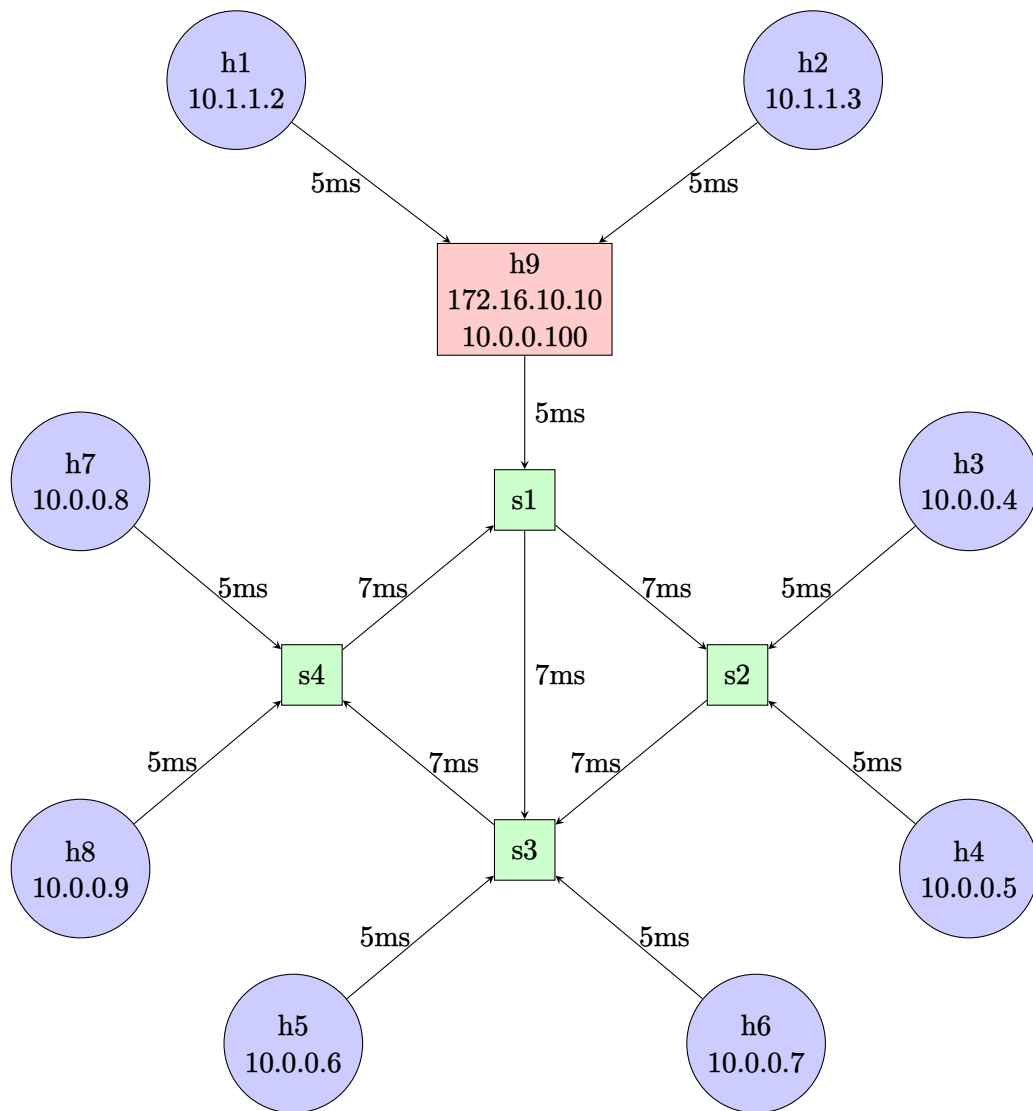
This experiment highlighted the importance of managing loops in switch-based networks. Without STP, loops caused broadcast storms and disrupted communication. Enabling STP ensured



a loop-free topology, stabilized MAC learning, and restored reliable connectivity. Thus, STP is crucial for maintaining robust and efficient layer-2 network operations.

Chapter 2

Configure Host-based NAT





```
mininet> net
h1 h1-eth0:h9-eth1
h2 h2-eth0:h9-eth2
h3 h3-eth0:s2-eth1
h4 h4-eth0:s2-eth2
h5 h5-eth0:s3-eth1
h6 h6-eth0:s3-eth2
h7 h7-eth0:s4-eth1
h8 h8-eth0:s4-eth2
h9 h9-eth0:s1-eth1 h9-eth1:h1-eth0 h9-eth2:h2-eth0
s1 lo: s1-eth1:h9-eth0 s1-eth2:s2-eth3 s1-eth3:s4-eth4 s1-eth4:s3-eth5
s2 lo: s2-eth1:h3-eth0 s2-eth2:h4-eth0 s2-eth3:s1-eth2 s2-eth4:s3-eth3
s3 lo: s3-eth1:h5-eth0 s3-eth2:h6-eth0 s3-eth3:s2-eth4 s3-eth4:s4-eth3 s3-eth5:s1-eth4
s4 lo: s4-eth1:h7-eth0 s4-eth2:h8-eth0 s4-eth3:s3-eth4 s4-eth4:s1-eth3
c0
```

Figure 2.1: The initial interfaces in the topology.

2.1 Topology and Changes

The original topology featured four Layer 2 switches (S1, S2, S3, S4) with hosts H1-H8 connected directly to switches. H9 was intended as a NAT host. Modifications included:

- Moved H1 and H2 from S1 to H9 with 5ms delay.
- Added H9-S1 link with 5ms delay.
- Configured H9 with internal IP 10.1.1.1/24 (**bridge br0**) and external IP 172.16.10.10/24.

2.2 Observations

2.2.1 Parts A and B Communication

- **Part A (Internal to External):**
 - Ping from H1 to H5 and H2 to H3 initially failed due to the absence of a default gateway and NAT configuration on H9.
- **Part B (External to Internal):**
 - Pings from H8 to H1 and H6 to H2 initially failed due to lack of DNAT and routing to the internal network (10.1.1.0/24).



```
[mininet> h1 ping -c 3 h5
ping: connect: Network is unreachable
[mininet> h2 ping -c 3 h3
ping: connect: Network is unreachable
[mininet> h8 ping -c 3 h1
ping: connect: Network is unreachable
[mininet> h6 ping -c 3 h2
ping: connect: Network is unreachable
```

Figure 2.2: Initially no connection is established across internal or external hosts of respective subnets.

2.2.2 H1 and H2 Communication

- H1 and H2 were not communicating initially because they lacked a common interface and were directly connected to S1 instead of H9. The addition of a bridge interface on H9 resolved this, enabling direct communication within the 10.1.1.0/24 subnet without requiring NAT.

```
[mininet> h1 ping -c 3 h2
PING 10.1.1.3 (10.1.1.3) 56(84) bytes of data.
64 bytes from 10.1.1.3: icmp_seq=1 ttl=64 time=45.7 ms
64 bytes from 10.1.1.3: icmp_seq=2 ttl=64 time=23.8 ms
64 bytes from 10.1.1.3: icmp_seq=3 ttl=64 time=23.8 ms

--- 10.1.1.3 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2004ms
rtt min/avg/max/mdev = 23.799/31.106/45.691/10.313 ms
```

Figure 2.3: h1 and h2 now become reachable owing to the bridge interface.

2.3 Issues and Rectifications

2.3.1 Initial Issues

- **Lack of Internal Connectivity:** No bridge or common interface existed for H1 and H2.
- **No Gateway for Internal Hosts:** H1 and H2 lacked default routes to reach external networks.
- **No External Access to Internal Network:** External hosts (H3-H8) could not route to 10.1.1.0/24 or perform reverse NAT.
- **Missing NAT Configuration:** IP forwarding and iptables rules were incomplete.



2.3.2 Rectifications with Code

- **Bridge Interface:**

- **Issue:** No common interface for H1 and H2 to communicate.
- **Rectification:** Created bridge **br0** on H9 to connect H1 and H2.
- **Code:**

```
1 h1, h2, h9 = net.get("h1", "h2", "h9")
2 h9.cmd("ip link add br0 type bridge")
3 h9.cmd("ip link set br0 up")
4 h9.cmd("ip link set h9-eth1 master br0") # h1-h9
5 h9.cmd("ip link set h9-eth2 master br0") # h2-h9
```

- **Explanation:** The **bridge (br0)** acts as a virtual switch, connecting H1 and H2 via H9's interfaces (h9-eth1, h9-eth2), allowing Layer 2 communication within 10.1.1.0/24.

- **Internal IP for NAT:**

- **Issue:** H9 lacked an internal IP to serve as a gateway.
- **Rectification:** Assigned 10.1.1.1/24 to **br0**.
- **Code:**

```
1 h9.cmd("ip addr add 10.1.1.1/24 dev br0") # IP for internal side
```

- **Explanation:** This IP serves as the default gateway for H1 (10.1.1.2) and H2 (10.1.1.3), enabling them to route traffic through H9.

- **Additional IP for Routing:**

- **Issue:** External hosts needed a route to H9's external network.
- **Rectification:** Added 10.0.0.100/24 to **h9-eth0**.
- **Code:**

```
1 h9.cmd("ip addr add 10.0.0.100/24 dev h9-eth0")
```

- **Explanation:** This secondary IP on h9-eth0 allows external hosts in the 10.0.0.0/24 network to route traffic to H9, facilitating communication with the internal network.

- **Default Routes:**

- **Issue:** H1 and H2 had no gateway to external networks.
- **Rectification:** Set default routes to 10.1.1.1.
- **Code:**

```
1 h1.cmd("ip route add default via 10.1.1.1")
2 h2.cmd("ip route add default via 10.1.1.1")
```

- **Explanation:** These routes direct all unknown traffic from H1 and H2 to H9 (10.1.1.1), which then handles NAT to external networks.



- **IP Forwarding:**

- **Issue:** H9 could not forward packets between interfaces.
- **Rectification:** Enabled IP forwarding.
- **Code:**

```
1 h9.cmd("sysctl -w net.ipv4.ip_forward=1")
```

- **Explanation:** This enables H9 to act as a router, forwarding packets between the internal bridge (br0) and external interface (h9-eth0).

- **IPTables:**

- **Issue:** No NAT rules existed for internal-to-external or external-to-internal traffic.
- **Rectification:** Added MASQUERADE and forwarding rules.
- **Code:**

```
1 h9.cmd("iptables -t nat -A POSTROUTING -s 10.1.1.0/24 -o h9-eth0  
-j MASQUERADE")
```

- **Explanation:** The POSTROUTING rule translates H1/H2's private IPs (10.1.1.0/24) to H9's public IP (172.16.10.10) for external access.

```
Chain POSTROUTING (policy ACCEPT 15 packets, 1024 bytes)
pkts bytes target      prot opt in      out     source         destination
 12   864 MASQUERADE    0    --  *      h9-eth0 10.1.1.0/24    0.0.0.0/0
```

Figure 2.4: POSTROUTING configuration added in NAT's IP Table

- **Route Add for External Hosts:**

- **Issue:** External hosts could not reach H9 or the internal network.
- **Rectification:** Added routes to 172.16.10.10 and 10.1.1.0/24 via 10.0.0.100.
- **Code:**

```
1 for host in [h3, h4, h5, h6, h7, h8]:
2     host.cmd("ip route add 172.16.10.10 via 10.0.0.100") # Route
   to NAT host
3     host.cmd("ip route add 10.1.1.0/24 via 10.0.0.100") # Route
   to internal network
```

- **Explanation:** These routes allow H3-H8 to send traffic to H9's external IP and the internal subnet, enabling reverse communication.

- **Enable Spanning Tree Protocol:**

- **Issue:** External hosts will be unreachable due to redundant links present in the topology.
- **Rectification:** Enable STP same as did in Question 1.



– Code:

```
1 for switch in ['s1', 's2', 's3', 's4']:  
2     net.get(switch).cmd('ovs-vsctl set bridge {} stp_enable=true'  
3     .format(switch))  
4     print("Waiting for STP to converge...")  
5     sleep(30)
```

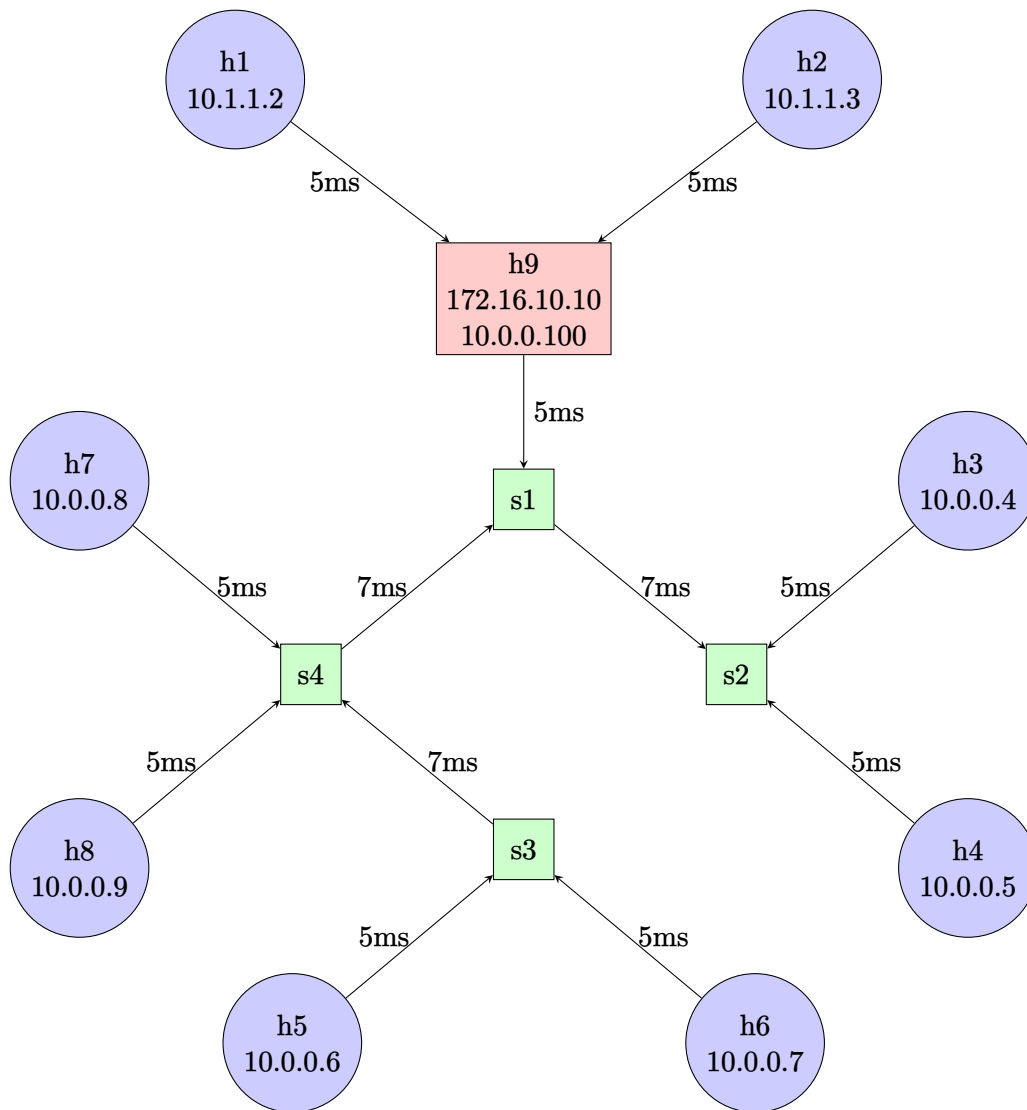


Figure 2.5: The links **s1-s3** and **s2-s3** have been disconnected as a result of blocking of redundant interfaces **s1-eth4** and **s2-eth4** by STP. Their identification is done as done in Ques 1.

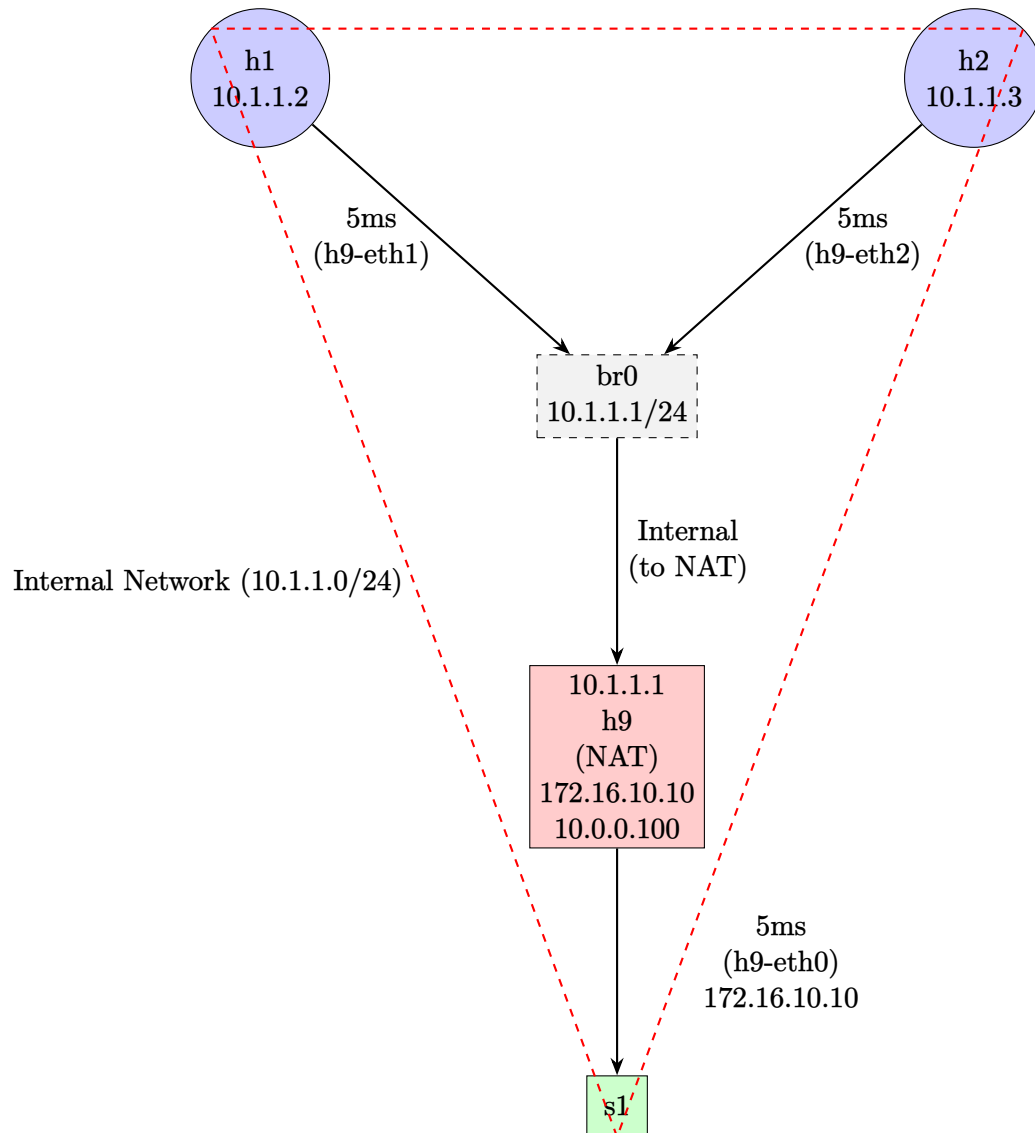


Figure 2.6: Figure showing NAT having aliases for both internal subnet (10.1.1./24) and to the external subnet (10.0.0./24) with the internal **bridge br0** interface



```
[mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3 h4 h5 h6 h7 h8 h9
h2 -> h1 h3 h4 h5 h6 h7 h8 h9
h3 -> h1 h2 h4 h5 h6 h7 h8 h9
h4 -> h1 h2 h3 h5 h6 h7 h8 h9
h5 -> h1 h2 h3 h4 h6 h7 h8 h9
h6 -> h1 h2 h3 h4 h5 h7 h8 h9
h7 -> h1 h2 h3 h4 h5 h6 h8 h9
h8 -> h1 h2 h3 h4 h5 h6 h7 h9
h9 -> h1 h2 h3 h4 h5 h6 h7 h8
*** Results: 0% dropped (72/72 received)
```

Figure 2.7: All hosts being reachable to each other.

The figure shows a Wireshark packet capture on interface s3-eth2. The packet list shows an ICMP Echo (ping) request from 10.0.0.7 to 10.0.0.7. The packet details show the Ethernet II frame, Internet Protocol Version 4, and Internet Control Message Protocol. The packet bytes show the ICMP echo request structure.

```
0000  3e 92 4a e5 87 f2 6e ee 60 f9 6f e1 08 00 45 00  >J...n...o...E...
0010  00 54 92 98 00 00 00 01 d3 a6 0a 00 00 07 0a 00  .T...@... ..
0020  00 64 00 00 5e 73 1e 71 00 11 6b d1 fa 67 00 00  .d...s...q...k...g...
0030  00 00 59 fe 04 00 00 00 00 00 11 12 13 14 15  .Y... ..
0040  16 17 18 19 1a 1b 1c 1d 1e 1f 20 21 22 23 24 25  .....0...!...%$%
0050  26 27 28 29 2a 2b 2c 2d 2e 2f 30 31 32 33 34 35  6'()*+,.../012345
0060  36 37 67
```

The packet capture shows a ping from h6 to h1. The packet list shows the ICMP echo request from 10.0.0.7 to 10.0.0.7. The packet details show the Ethernet II frame, Internet Protocol Version 4, and Internet Control Message Protocol. The packet bytes show the ICMP echo request structure.

```
iperf Done.
mininet> h1 ping h6
PING 10.0.0.7 (10.0.0.7) 56(84) bytes of data:
64 bytes from 10.0.0.7: icmp_seq=1 ttl=63 time=68.3 ms
64 bytes from 10.0.0.7: icmp_seq=2 ttl=63 time=60.3 ms
64 bytes from 10.0.0.7: icmp_seq=3 ttl=63 time=62.6 ms
64 bytes from 10.0.0.7: icmp_seq=4 ttl=63 time=58.2 ms
64 bytes from 10.0.0.7: icmp_seq=5 ttl=63 time=55.3 ms
64 bytes from 10.0.0.7: icmp_seq=6 ttl=63 time=58.2 ms
64 bytes from 10.0.0.7: icmp_seq=7 ttl=63 time=48.9 ms
64 bytes from 10.0.0.7: icmp_seq=8 ttl=63 time=51.7 ms
64 bytes from 10.0.0.7: icmp_seq=9 ttl=63 time=53.1 ms
64 bytes from 10.0.0.7: icmp_seq=10 ttl=63 time=53.7 ms
64 bytes from 10.0.0.7: icmp_seq=11 ttl=63 time=52.3 ms
```

Figure 2.8: **h1 ping h6**: wireshark packet capture shows h6 (10.0.0.7) reaches the NAT's same subnet (10.0.0.100) rather than h1 directly.



2.4 Finally all the Communication Tests

2.4.1 Communication to an external host from an internal host

1. Ping to h5 from h1: (Expected RTT = 58 ms)
2. Ping to h3 from h2: (Expected RTT = 44 ms)

Table 2.1: Average RTTs for 30s Ping Tests After STP Convergence

Ping Test	Attempt 1 (ms)	Attempt 2 (ms)	Attempt 3 (ms)	Avg. Time (ms)
h1 → h5	60.677	58.969	58.746	59.464
h2 → h3	46.068	44.703	45.023	45.265

2.4.2 Test communication to an internal host from an external host

1. Ping to h1 from h8: (Expected RTT = 44 ms)
2. Ping to h2 from h6: (Expected RTT = 58 ms)

Table 2.2: Average RTTs for 30s Ping Tests After STP Convergence

Ping Test	Attempt 1 (ms)	Attempt 2 (ms)	Attempt 3 (ms)	Avg. Time (ms)
h8 → h1	46.005	44.980	44.350	45.112
h6 → h2	60.786	58.717	58.814	59.439

2.4.3 iperf tests: 3 tests of 120s each.

1. Run iperf3 server in h1 and iperf3 client in h6 at port 1212.

```
1 h1 iperf3 -s -p 1212 -D
2 h6 iperf3 -c h1 -p 1212 -t 120
```

2. Run iperf3 server in h8 and iperf3 client in h2 at port 1212.

```
1 h8 iperf3 -s -p 1212 -D
2 h2 iperf3 -c h8 -p 1212 -t 120
```



```
h5 h5-eth0:s3-eth1
h6 h6-eth0:s3-eth2
h7 h7-eth0:s4-eth1
h8 h8-eth0:s4-eth2
h9 h9-eth0:s1-eth1 h9-eth1:h1-eth0 h9-eth2:h2-eth0
s1 lo: s1-eth1:h9-eth0 s1-eth2:s2-eth3 s1-eth3:s4-eth4 s1-eth4:s3-eth5
s2 lo: s2-eth1:h3-eth0 s2-eth2:h4-eth0 s2-eth3:s1-eth2 s2-eth4:s3-eth3
s3 lo: s3-eth1:h5-eth0 s3-eth2:h6-eth0 s3-eth3:s2-eth4 s3-eth4:s4-eth3 s3-eth5:s1-eth4
s4 lo: s4-eth1:h7-eth0 s4-eth2:h8-eth0 s4-eth3:s3-eth4 s4-eth4:s1-eth3
c0
mininet> h6 iperf3 -c h1 -p 1212 -t 120
Connecting to host 10.1.1.2, port 1212
[ 5] local 10.0.0.7 port 43586 connected to 10.1.1.2 port 1212
[ ID] Interval           Transfer     Bitrate      Retr  Cwnd
[ 5] 0.00-1.00 sec       16.5 MBytes  138 Mbits/sec    0   1.28 MBytes
[ 5] 1.00-2.00 sec       37.5 MBytes  315 Mbits/sec    0   2.87 MBytes
[ 5] 2.00-3.00 sec       60.0 MBytes  502 Mbits/sec    0   5.91 MBytes
[ 5] 3.00-4.01 sec       97.5 MBytes  815 Mbits/sec   406   8.26 MBytes
[ 5] 4.01-5.00 sec      104 MBytes  876 Mbits/sec    0   8.26 MBytes
[ 5] 5.00-6.01 sec      96.2 MBytes  801 Mbits/sec    0   8.26 MBytes
[ 5] 6.01-7.00 sec      97.5 MBytes  823 Mbits/sec    0   8.26 MBytes
[ 5] 7.00-8.00 sec      96.2 MBytes  807 Mbits/sec    0   8.26 MBytes
[ 5] 8.00-9.01 sec      90.0 MBytes  750 Mbits/sec    0   8.26 MBytes
[ 5] 9.01-10.01 sec     90.0 MBytes  752 Mbits/sec    0   8.26 MBytes
```

Figure 2.9: Server in h1 and iperf3 Client in h6. Also, see packets in the wireshark.

```
69365 126.723788691 10.1.1.2 10.0.0.7 TCP 54 1212 -> 43586 [RST] Seq=2 Win=0 Len=0
69366 126.743849657 10.1.1.2 10.0.0.7 TCP 66 1212 -> 43584 [ACK] Seq=6 Ack=148 Win=43520 Len=0 TSval=302555..
69367 126.749711865 10.0.0.7 10.1.1.2 TCP 359 43584 -> 1212 [PSH, ACK] Seq=148 Ack=6 Win=42496 Len=293 TSval..
69368 126.75153109 10.0.0.7 10.1.1.2 TCP 66 1212 -> 43584 [ACK] Seq=6 Ack=119 Win=43520 Len=0 TSval=302555..
69369 126.794584687 10.1.1.2 10.0.0.7 TCP 70 1212 -> 43584 [PSH, ACK] Seq=6 Ack=441 Win=43520 Len=4 TSval=3..
69370 126.800427898 10.0.0.7 10.1.1.2 TCP 66 43584 -> 1212 [ACK] Seq=441 Ack=10 Win=42496 Len=0 TSval=37388..

Frame 69368: 66 bytes on wire (528 bits), 66 bytes captured (528 bits) on interface s3-eth2, id 0
Ethernet II, Src: 3e:92:4a:e5:87:f2 (3e:92:4a:e5:87:f2), Dst: 6e:ee:60:f9:6f:e1 (6e:ee:60:f9:6f:e1)
Internet Protocol Version 4, Src: 10.1.1.2, Dst: 10.0.0.7
Transmission Control Protocol, Src Port: 1212, Dst Port: 43584, Seq: 6, Ack: 441, Len: 0

0000  6e ee 60 f9 6f e1 3e 92 4a e5 87 f2 08 00 45 00  n...o...J....E
0010  00 34 3f bf 40 00 3f 06 e6 fb 0a 01 01 02 0a 00  .4?@?.....
0020  00 07 04 bc aa 40 09 e4 b6 ed a2 ea 03 b2 80 10  ....@.....
0030  00 55 15 30 00 00 01 01 08 0a 12 08 9f e4 0e 0e  .U0.....
0040  0b 85  .

s3-eth2: <live capture in progress> Packets: 69412 · Displayed: 69412 (100.0%) Profile: Default

[ 5] 110.00-111.00 sec 75.0 MBytes 629 Mbits/sec 0 8.26 MBytes
[ 5] 111.00-112.00 sec 88.8 MBytes 748 Mbits/sec 0 8.26 MBytes
[ 5] 112.00-113.00 sec 95.0 MBytes 796 Mbits/sec 0 8.26 MBytes
[ 5] 113.00-114.00 sec 102 MBytes 861 Mbits/sec 0 8.26 MBytes
[ 5] 114.00-115.00 sec 114 MBytes 953 Mbits/sec 0 8.26 MBytes
[ 5] 115.00-116.00 sec 131 MBytes 1.10 Gbits/sec 0 8.26 MBytes
[ 5] 116.00-117.00 sec 132 MBytes 1.11 Gbits/sec 0 8.26 MBytes
[ 5] 117.00-118.00 sec 138 MBytes 1.15 Gbits/sec 0 8.26 MBytes
[ 5] 118.00-119.00 sec 139 MBytes 1.16 Gbits/sec 0 8.26 MBytes
[ 5] 119.00-120.00 sec 139 MBytes 1.16 Gbits/sec 0 8.26 MBytes

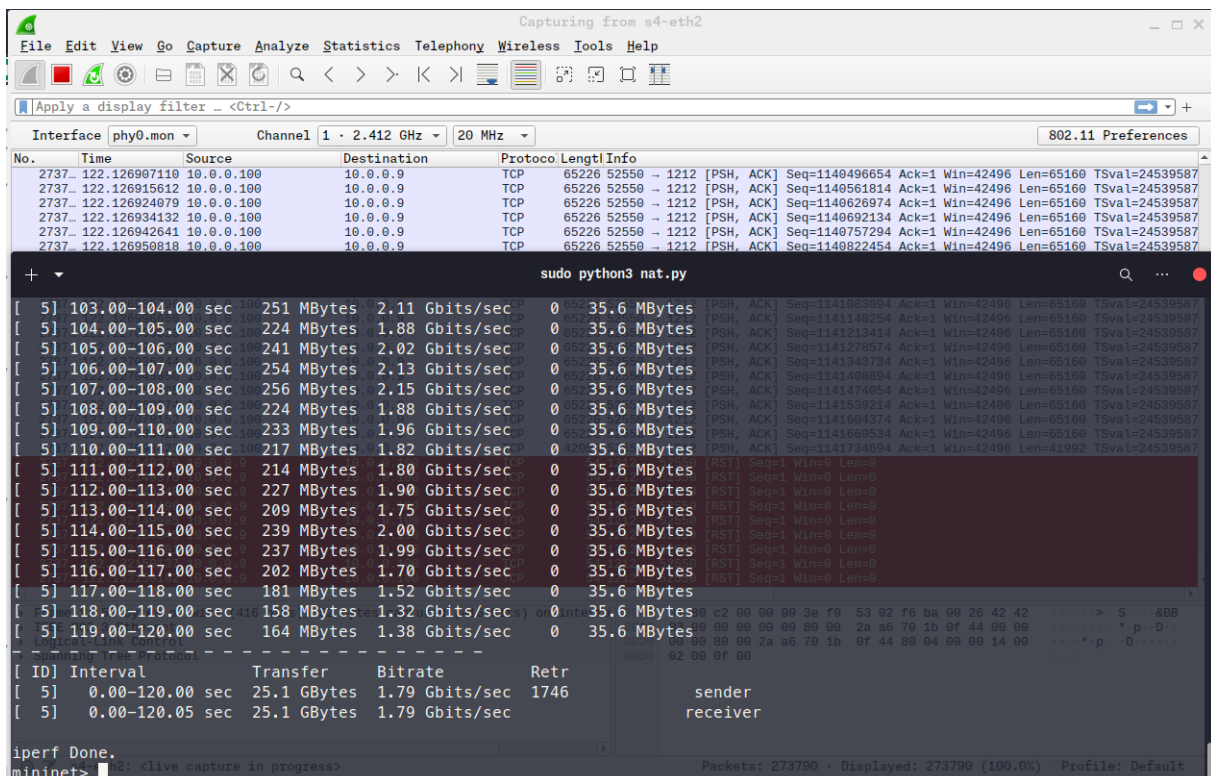
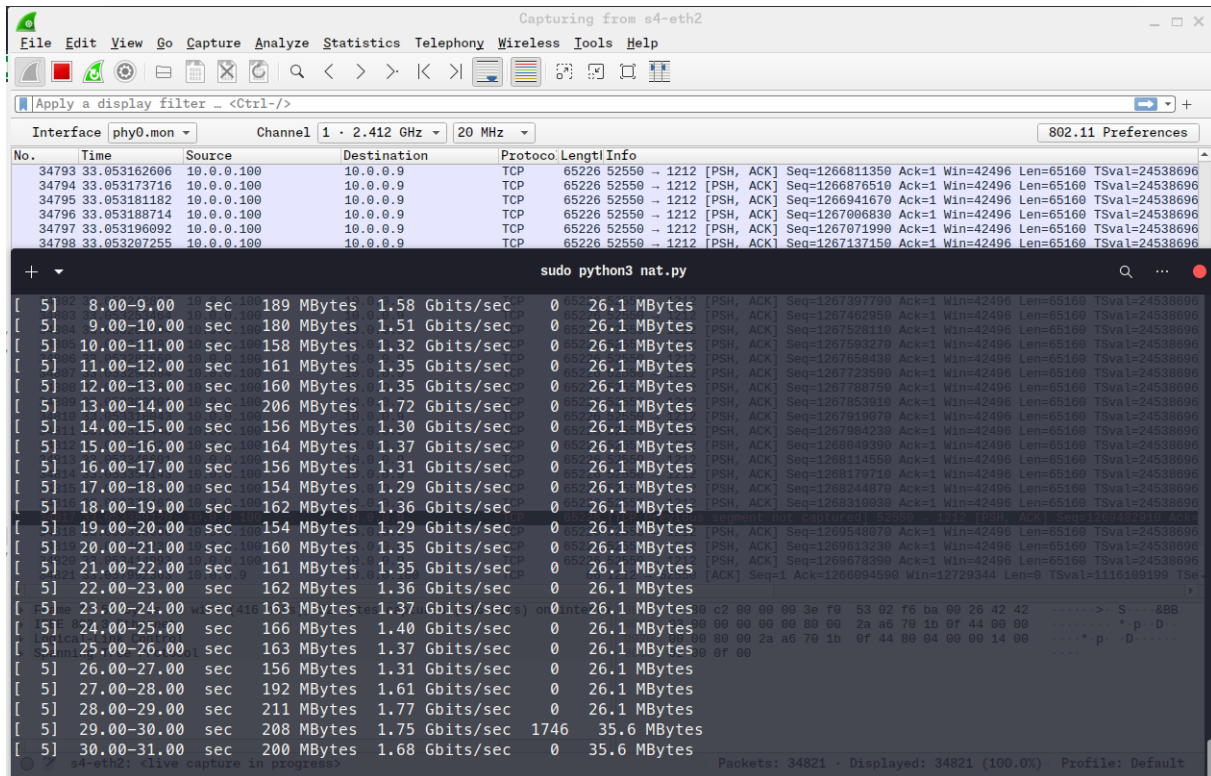
[ ID] Interval           Transfer     Bitrate      Retr
[ 5] 0.00-120.00 sec 13.1 GBytes 938 Mbits/sec 537
[ 5] 0.00-120.08 sec 13.1 GBytes 938 Mbits/sec

iperf Done.
mininet>
```

Figure 2.10: Final Result of running iperf: can see the completed packets in wireshark.



Chapter 2 : Configure Host-based NAT



Chapter 3

Distributed Asynchronous Distance Vector Routing

3.1 Distance Vector Routing Algorithm

The [Distance Vector Routing Algorithm](#), also known as the Bellman-Ford algorithm, operates on the principle that each node maintains a vector of distances to all other nodes in the network and periodically shares this vector with its directly connected neighbors. The key steps are:

1. **Initialization:** Each node initializes its distance table with the costs to its directly connected neighbors and sets the cost to non-neighbors as infinity (e.g., 999). The distance to itself is zero.
2. **Information Exchange:** Nodes send their distance vectors to neighbors, containing the minimum cost to reach each node.
3. **Update:** Upon receiving a distance vector from a neighbor, a node updates its distance table using the Bellman-Ford equation:

$$D_x(y) = \min_v [c(x, v) + D_v(y)]$$

where $D_x(y)$ is the cost from node x to node y , $c(x, v)$ is the cost of the link from x to neighbor v , and $D_v(y)$ is the cost reported by v to y .

4. **Convergence:** If the minimum cost to any destination changes, the node sends an updated distance vector to its neighbors, continuing until all tables stabilize.

This process is asynchronous, meaning updates occur independently as packets are received, and the network emulator ensures in-order delivery without loss.

3.2 Implementation Approach

The implementation involves [four source files](#) (`node0.c`, `node1.c`, `node2.c`, `node3.c`), each corresponding to a node, and a main emulator file (`distance_vector.c`). The approach includes:

3.2.1 Data Structures

- **Distance Table:** A 4x4 matrix (`costs[i][j]`) where `costs[i][j]` is the cost from node i to node j via neighbor j . Initialized with direct costs and infinity for non-neighbors.
- **Minimum Cost Array:** A vector (`min_costsX`) storing the minimum cost to each node, updated after table changes.



- **Routing Packet:** A struct `rtpkt` containing `sourceid`, `destid`, and `mincost[4]` fields, used to exchange distance vectors.

3.2.2 Code Organization

Each node file implements:

- **Initialization Function (`rtinitX`):** Sets up the distance table with direct costs (e.g., `wts0 = {0, 1, 3, 7}` for node 0) and sends the initial distance vector to neighbors.
- **Update Function (`rtupdateX`):** Updates the distance table based on received packets, recalculates minimum costs, and sends updates if changes occur.
- **Helper Functions:** `make_distance_vectorX` computes the minimum costs, and `send_packetX` sends the distance vector to neighbors.

3.2.3 Key Logic

- **Initialization:** For node 0, `dt0.costs[1][1] = 1`, `dt0.costs[2][2] = 3`, `dt0.costs[3][3] = 7`, with others set to 999.
- **Update:** For a packet from node `s`, `dtX.costs[i][s] = wtsX[s] + rcvdpkt->mincost[i]`, and the minimum cost is recomputed. If it changes, neighbors are notified.
- **Error Handling:** The code uses `INF (999)` for unreachable nodes and includes tracing via `TRACE`.

3.2.4 Refinements

- Removed global `src_node` to avoid duplicate symbol errors, hardcoding node IDs in `send_packetX`.
- Used local `src_node` in `rtupdateX` for the packet source ID.
- Made a global `make_distance_vector()` function for each of the node modules.
- Created a `send_packetX()` for each of the modules to send the packets to only the neighbours.

Listing 3.1: Example: `rtinit0` in `node0.c`

```
1 void rtinit0() {
2     for (int i = 0; i < 4; i++) {
3         for (int j = 0; j < 4; j++) {
4             if (i == j) dt0.costs[i][j] = wts0[i];
5             else dt0.costs[i][j] = INF;
6         }
7     }
8     if (TRACE > 0) printdt0(&dt0);
9     make_distance_vector0();
10    send_packet0();
11 }
```




Listing 3.2: `make_distance_vector0` and `send_packet0()` in `node0.c`

```
1 int min_macro0(int a, int b)
2 {
3     return (a < b) ? a : b;
4 }
5
6 int min_cost0(int* a)
7 {
8     return min_macro0(min_macro0(a[0], a[1]), min_macro0(a[2], a[3]));
9 }
10
11 // Send the distance vector to all directly connected neighbors
12 // min_costs[i] = min_cost(dt0.costs[i]);
13 // i.e. For destinations marked on columns, minimum cost to reach them
14 // from node 0 is the minimum of the costs to reach them via all other
   nodes
15 int min_costs0[4];
16 void make_distance_vector0()
17 {
18     for (int i = 0; i < 4; i++)
19         min_costs0[i] = min_cost0(dt0.costs[i]);
20 }
21
22 void send_packet0()
23 {
24     struct rtpkt packet;
25     packet.sourceid = 0;
26     for (int j = 0; j < 4; j++)
27         packet.mincost[j] = min_costs0[j];
28
29     for (int i = 0; i < 4; i++)
30     {
31         if (i == 0) continue;
32         packet.destid = i;
33         tolayer2(packet);
34     }
35 }
```

3.3 Role of `distance_vector.c`

The `distance_vector.c` file serves as the network emulator and main routine, which students should not modify. Its key functions include:

- **Initialization (`init`):** Calls `rtinitX` for all nodes and sets up the event list for link changes (disabled with `LINKCHANGES=0`).
- **Event Handling (`main`):** Processes events (e.g., packet arrivals from `FROM_LAYER2`) by invoking `rtupdateX` and handles link changes via `linkhandlerX`.
- **Packet Transmission (`tolayer2`):** Simulates packet delivery to neighbors with variable delay (1-10 time units), checking connectivity via a `connectcosts` matrix.
- **Utilities:** Provides `jimsrand()` for random delays and `insertevent()` for event scheduling.



Listing 3.3: rtupdate0 in node0.c

```
1 void rtupdate0(struct rtpkt *rcvdpkt)
2 {
3     int src = rcvdpkt->sourceid;
4
5     int changed = 0;
6     int old_min_costs[4];
7     for (int i = 0; i < 4; i++)
8         old_min_costs[i] = min_costs0[i];
9
10    // Update the distance table for node 0
11    // i.e. if node 0's own minimum cost to another node changes as a
12    // result of the update
13    // node 0 informs its directly connected neighbors of this change in
14    // minimum cost
15    // hence node 0's own minimum cost to another node src_node = dt0.
16    // costs[src_node][src_node] + rcvdpkt->mincost[i]
17
18    for (int i = 0; i < 4; i++)
19    {
20        int new_cost = dt0.costs[src][src] + rcvdpkt->mincost[i];
21        if (new_cost < INF)
22            dt0.costs[i][src] = new_cost;
23        else
24            dt0.costs[i][src] = INF;
25    }
26
27    if (TRACE > 0) printdt0(&dt0);
28
29    make_distance_vector0();
30    for (int i = 0; i < 4; i++)
31    {
32        if (min_costs0[i] != old_min_costs[i])
33        {
34            changed = 1;
35            break;
36        }
37    }
38
39    if (changed)
40        send_packet0();
41 }
```



```
MAIN: rcv event, t=11.022, at 2 src: 0, dest: 2, contents:  0  1  2  7
      via
      D2 |  0  1  3
      ----|-----
      0 |  3  2  7
dest 1 |  4  1  5
      3 | 10  9  2
MAIN: rcv event, t=12.792, at 2 src: 1, dest: 2, contents:  1  0  1  3
      via
      D2 |  0  1  3
      ----|-----
      0 |  3  2  7
dest 1 |  4  1  5
      3 | 10  4  2
MAIN: rcv event, t=13.977, at 2 src: 0, dest: 2, contents:  0  1  2  5
      via
      D2 |  0  1  3
      ----|-----
      0 |  3  2  7
dest 1 |  4  1  5
      3 |  8  4  2
MAIN: rcv event, t=15.855, at 2 src: 3, dest: 2, contents:  4  3  2  0
      via
      D2 |  0  1  3
      ----|-----
      0 |  3  2  6
dest 1 |  4  1  5
      3 |  8  4  2
MAIN: rcv event, t=17.443, at 2 src: 0, dest: 2, contents:  0  1  2  4
      via
      D2 |  0  1  3
      ----|-----
      0 |  3  2  6
dest 1 |  4  1  5
      3 |  7  4  2

Simulator terminated at t=17.443218, no packets in medium
(base) guntas13 JetBrains Projects/routing-nat-loop master
```

Figure 3.1: Running the routing algorithm.